

```
from sklearn.preprocessing import LabelBinarizer
from sklearn.metrics import classification_report
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.datasets import mnist

import matplotlib.pyplot as plt
import numpy as np

print("[INFO] accessing MNIST...")

[INFO] accessing MNIST...

((trainX, trainY), (testX, testY)) = mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 0s 0us/step

trainX = trainX.reshape((trainX.shape[0], 28 * 28 * 1))
testX = testX.reshape((testX.shape[0], 28 * 28 * 1))

trainX = trainX.astype("float32") / 255.0
testX = testX.astype("float32") / 255.0

lb = LabelBinarizer()
trainY = lb.fit_transform(trainY)
testY = lb.transform(testY)

model = Sequential()
model.add(Dense(256, input_shape=(784,), activation="relu"))
model.add(Dense(128, activation="relu"))
model.add(Dense(64, activation="relu"))
model.add(Dense(10, activation="softmax"))

print("[INFO] training network...")
Adm = Adam(0.01)
model.compile(loss="categorical_crossentropy", optimizer=Adm,
metrics=["accuracy"])
H = model.fit(trainX, trainY, validation_data=(testX, testY),
epochs=100, batch_size=128)
```

```

epoch 90/100
469/469 [=====] - 4s 8ms/step - loss: 0.0367 - accuracy: 0.9944 - val_loss: 0.2526 - val_accuracy: 0.975
Epoch 91/100
469/469 [=====] - 4s 8ms/step - loss: 0.0179 - accuracy: 0.9960 - val_loss: 0.2878 - val_accuracy: 0.977
Epoch 92/100
469/469 [=====] - 5s 10ms/step - loss: 0.0367 - accuracy: 0.9942 - val_loss: 0.3848 - val_accuracy: 0.97
Epoch 93/100
469/469 [=====] - 4s 8ms/step - loss: 0.0277 - accuracy: 0.9949 - val_loss: 0.3882 - val_accuracy: 0.975
Epoch 94/100
469/469 [=====] - 4s 8ms/step - loss: 0.0274 - accuracy: 0.9964 - val_loss: 0.3445 - val_accuracy: 0.977
Epoch 95/100
469/469 [=====] - 4s 9ms/step - loss: 0.0179 - accuracy: 0.9970 - val_loss: 0.4133 - val_accuracy: 0.976
Epoch 96/100
469/469 [=====] - 4s 8ms/step - loss: 0.0353 - accuracy: 0.9936 - val_loss: 0.3252 - val_accuracy: 0.976
Epoch 97/100
469/469 [=====] - 4s 8ms/step - loss: 0.0183 - accuracy: 0.9966 - val_loss: 0.3795 - val_accuracy: 0.974
Epoch 98/100
469/469 [=====] - 4s 10ms/step - loss: 0.0194 - accuracy: 0.9964 - val_loss: 0.3829 - val_accuracy: 0.97
Epoch 99/100
469/469 [=====] - 4s 8ms/step - loss: 0.0172 - accuracy: 0.9965 - val_loss: 0.4113 - val_accuracy: 0.972
Epoch 100/100

```

```

print("[INFO] evaluating network...")
predictions = model.predict(testX, batch_size=128)
print(classification_report(testY.argmax(axis=1),
    predictions.argmax(axis=1),
    target_names=[str(x) for x in lb.classes_]))

```

```

[INFO] evaluating network...
79/79 [=====] - 0s 4ms/step

```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	980
1	0.99	0.99	0.99	1135
2	0.96	0.98	0.97	1032
3	0.96	0.98	0.97	1010
4	0.98	0.98	0.98	982
5	0.98	0.97	0.98	892
6	0.98	0.98	0.98	958
7	0.98	0.97	0.98	1028
8	0.97	0.97	0.97	974
9	0.98	0.97	0.98	1009
accuracy			0.98	10000
macro avg	0.98	0.98	0.98	10000
weighted avg	0.98	0.98	0.98	10000

```

plt.style.use("ggplot")
plt.figure()
plt.plot(np.arange(0, 100), H.history["loss"], label="train_loss")
plt.plot(np.arange(0, 100), H.history["val_loss"], label="val_loss")
plt.plot(np.arange(0, 100), H.history["accuracy"], label="train_acc")
plt.plot(np.arange(0, 100), H.history["val_accuracy"],
    label="val_acc")
plt.title("Training Loss and Accuracy")
plt.xlabel("Epoch #")
plt.ylabel("Loss/Accuracy")
plt.legend()

```

<matplotlib.legend.Legend at 0x7dd41427e350>

